

FlashPCA: fast sparse canonical correlation analysis of genomic data — supplementary material

Gad Abraham and Michael Inouye

April 4, 2016

1 Reproducibility

Code to reproduce these experiments is at <https://github.com/gabraham/scca-paper>.

2 HapMap data preprocessing and quality control

The HapMap3 phase III [1] genotypes were obtained from ftp://ftp.ncbi.nlm.nih.gov/hapmap/genotypes/2010-05_phaseIII/plink_format/. Gene expression levels were obtained from <http://www.ebi.ac.uk/arrayexpress/experiments/E-MTAB-264> and <http://www.ebi.ac.uk/arrayexpress/experiments/E-MTAB-198> (the CEU subset).

Within each of the eight populations (CEU, CHB, GIH, JPT, LWK, MEX, MKK, YRI), we excluded individuals who were non-founders, had genotyping missingness $>10\%$, or did not have matching gene expression data, resulting in 709 individuals in total. In addition, within each population, we excluded non-autosomal SNPs, SNPs with MAF $<1\%$, missingness $>10\%$, and deviation from Hardy-Weinberg equilibrium $P < 10^{-6}$ using PLINK 1.9 [2, 3]. Finally, we combined all eight population into one dataset, consisting of the post-QC autosomal SNPs within each dataset, a total of 973,983 SNPs.

In the analysis, any remaining missing genotypes were randomly imputed according to the frequencies of the non-missing observations (equivalent on average to mean imputation).

For the gene expression data, we used a subset consisting of the 21,800 probes that were analysed by [4], utilising the original authors' normalised data. Following [4], we performed PCA on the genotypes within each population, and for the GIH, MEX, MKK, and LWK regressed out 10 PCs of the genotypes (as well as intercept) from the corresponding gene expression levels, in order to adjust for the higher levels of admixture within these populations. We further filtered probes with low variance (std. dev. <0.1), leaving 18,379 probes. Both the gene expression levels and the genotypes were standardised to zero-mean and unit-variance.

3 Comparison of predictive power with simulated gene expression data

The number of samples in the HapMap3 data ($n = 709$) does not provide adequate statistical power to detect weak correlations or difference in correlations between two competing methods, particularly when cross-validation further reduces the sample size in the test data (3-fold cross-validation leads to a test set of ~ 236 individuals). For example, there is power of only 34% to detect a correlation $\rho = 0.1$ at an $\alpha = 0.05$ with 236 samples, and only 8% power to detect a difference in correlations $\Delta\rho = 0.05$ at $\alpha = 0.05$ (i.e., comparing whether two test-set correlations achieved by competing methods are significantly different from one another).

Hence, we chose to simulate gene expression data with strong associations with the genotypes, allowing for higher correlations to be observed and meaningfully compared. Utilising 10,000 SNPs from HapMap3

chromosome 1, we simulated 1,000 gene expression levels as

$$\mathbf{Y} = \mathbf{XB} + \mathbf{E},$$

where \mathbf{X} are the genotypes ($n \times p$ matrix, n individuals and p SNPs), \mathbf{B} is a $p \times m$ matrix of weights (effect sizes for each pair of SNP and phenotype), and is an $n \times m$ matrix representing the error (noise). To match the sparsity assumptions of SCCA, \mathbf{B} was chosen to be a mixture of weights $\{0.001, 1\}$ with proportions 0.9999 and 0.0001 (across all $n \times m$ entries), respectively. Note that a value of 0.001 was used rather than zero, in order to prevent some probes from having zero genetic variance. Each column $k = 1, \dots, m$ of \mathbf{E} was $E_k \sim \mathcal{N}(0, \frac{1-h^2}{h^2} \text{var}((\mathbf{XB})_k))$, and $h^2 = 0.1$.

Due to the different formulation used by `PMA::CCA` versus that of `flashpcaR::scca` (constrained versus penalised, respectively), it is not straightforward to compare the two models directly for a given penalisation level. Hence, we compared the best predictive performance achieved by the tools in cross-validation. We used 3-fold cross-validation over a 2D grid of 30×25 penalties, estimating one pair of canonical vectors. The final predictive power was computed as the average Pearson correlation $\bar{\rho}$ in the $k = 1, \dots, 3$ test folds:

$$\bar{\rho} = \frac{1}{3} \sum_{k=1}^3 \text{Cor}(\mathbf{X}_{\text{test}}^k u^k, \mathbf{Y}_{\text{test}}^k v^k).$$

The cross-validation folds were identical for both methods tested. The maximum of the average test Pearson correlation was identical for `flashpcaR::scca` and `PMA::CCA` ($\rho=0.936$), completing in 5m and 24m, respectively (parallelising over 3 cores).

4 Timing experiments

For timing of `flashpcaR::scca` and `PMA::CCA` we used contiguous subsets of HapMap3 chromosome 1 (1000, 5000, 10,000, 20,000, and 50,000 SNPs, out of 79,011 SNPs in total) and contiguous subsets of the 18,379 real gene expression probes (1000, 10,000, and all 18,379 probes) [4].

We used the R package `microbenchmark` [5] to run 30 replications of each timing experiment. For all experiments we estimated one pair of canonical vectors (u_1, v_1). For the results in the main text, we initialised (“warm started”) v_1 to a standard normally-distributed vector of variates $\sim \mathcal{N}(0, 1)$. `PMA::CCA` and `flashpcaR::scca` allow the user to provide their own initialisation¹, and we experimented with other forms, including using the column means of the gene expression data and the rank-1 singular value decomposition (SVD) $\mathbf{X}^T \mathbf{Y} \approx u_1 d_1 v_1^T$. The overall trend of `flashpcaR` being several-fold faster than `PMA` was consistent across all three initialisation methods (Figure 1).

All experiments were run in R 3.2.2 [6] (with the original LAPACK and BLAS libraries included in R) on 64-bit Ubuntu Linux 12.04 on an Intel Xeon CPU E7-4830 v2 @ 2.20GHz. Time for the commandline `flashpca` include loading of data into RAM and all preprocessing (e.g., standardising the variables). We used `flashpca` v1.2.6 (<https://github.com/gabraham/flashpca>) and `PMA` v1.0.9 [7]. For `PMA::CCA`, we increased the maximum number of iterations to match that used by `flashpcaR::scca` (default=1000), in order to prevent early termination of the algorithm before adequate numerical convergence was achieved.

5 Parallelising grid search for penalty optimisation

As described in Section 3, using multiple cores can speed up the penalty grid search for `PMA` and `flashpcaR`. Within R, this can be achieved using the `foreach` [8] and `doMC` [9] packages. We recommend using coarse-grain parallelisation for cross-validation, e.g., 5 cores for 5-fold cross-validation. Examples are given in the code at <https://github.com/gabraham/scca-paper>.

The commandline tool `flashpca` currently does not support built-in cross-validation (as of v1.2.6). We recommend splitting the data into training/test folds using PLINK and running `flashpca` on these subsets, possibly using GNU parallel [10].

¹The commandline version `flashpca` currently only supports random initialisation.




Figure 1: Timing (median of 30 runs) of SCCA implemented in the `flashpcaR` (R package) and in `flashpca` (stand-alone commandline tool), compared with SCCA from PMA, using subsets of the HapMap3 dataset with real gene expression levels as phenotypes. We compared three schemes for initialising v_1 : (i) “mean”: column means of the gene expression data; (ii) “rand”: normally-distributed variates $\mathcal{N}(0, 1)$; and (iii) “svd”: 1st right singular value of $\mathbf{X}^T \mathbf{Y}$. Note that the commandline `flashpca` currently only supports random initialisation.

References

- [1] International HapMap 3 Consortium. Integrating common and rare genetic variation in diverse human populations. *Nature*, 467:52–58, 2010.
- [2] S. Purcell, B. Neale, K. Todd-Brown, L. Thomas, M. A. Ferreira, D. Bender, J. Maller, P. Sklar, P. I. de Bakker, M. J. Daly, and P. C. Sham. PLINK: a tool set for whole-genome association and population-based linkage analyses. *Am. J. Hum. Genet.*, 81:559–575, 2007.
- [3] C Chang, C Chow, L Tellier, S Vattikuti, S Purcell, and J Lee. Second-generation plink: rising to the challenge of larger and richer datasets. *GigaScience*, 4(1):7, 2015. ISSN 2047-217X. doi: 10.1186/s13742-015-0047-8.
- [4] B. Stranger et al. Patterns of Cis Regulatory Variation in Diverse Human Populations. *PLoS Genet*, 8(4):e1002639, 2012. doi: 10.1371/journal.pgen.1002639.
- [5] O Mersmann. *microbenchmark: Accurate Timing Functions*, 2015. URL <http://CRAN.R-project.org/package=microbenchmark>. R package version 1.4-2.1.
- [6] R Development Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2015. URL <http://www.R-project.org>. ISBN 3-900051-07-0.
- [7] D Witten, R Tibshirani, S Gross, and B Narasimhan. *PMA: Penalized Multivariate Analysis*, 2013. URL <http://CRAN.R-project.org/package=PMA>. R package version 1.0.9.
- [8] Revolution Analytics and Steve Weston. *foreach: Provides Foreach Looping Construct for R*, 2015. URL <http://CRAN.R-project.org/package=foreach>. R package version 1.4.3.
- [9] Revolution Analytics and Steve Weston. *doMC: Foreach Parallel Adaptor for 'parallel'*, 2015. URL <http://CRAN.R-project.org/package=doMC>. R package version 1.3.4.
- [10] O. Tange. GNU Parallel - The Command-Line Power Tool. *login: The USENIX Magazine*, 36(1):42–47, Feb 2011. doi: 10.5281/zenodo.16303. URL <http://www.gnu.org/s/parallel>.